

Single Source Plagiarism Detection in Intro Level Computer Science Classes Using Keystroke Metadata

Jesse McDonald
University of Hawai'i at Mānoa
jamcd@hawaii.edu

Scott Robertson
University of Hawai'i at Mānoa
scottpr@hawaii.edu

Abstract

Plagiarism in intro level computer science courses is a problem. The ability for students to pass classes by copying assignments without learning the material is hurting the industry as a whole. Previously available tools for combating plagiarism are failing as online resources increases in quality and availability. These can range from publicly accessible solutions, to pay-walled solutions, and even to bespoke custom made solutions individual to each student. Freely available generative AI capable of solving intro level problems further exacerbates this problem. Thus we propose a novel approach where plagiarism is determined from metadata embedded by a specialized IDE into student submissions. We find this approach is generally more reliable than traditional plagiarism detection methods in both theory and practice.

1. Introduction

Introductory Computer Science classes are often students first look at programming. In these classes people who don't know how to program, learn to program for the first time. Or that is the assumption. However, due to the nature of computer science it is often possible to pass an intro to computer science course without writing a single line of code by plagiarizing assignments.

1.1. Plagiarisms Logical Consequences

Plagiarism is particularly troublesome in intro classes. Since the intro classes are often a student's first experience programming, if a student passes an intro class by plagiarizing it is likely that student does not understand part or all of the course. Since higher programming courses assume you know the basics, these students have only 2 choices going forward. Firstly they can work hard and put in extra time to teach themselves the missed information. This is unlikely, as this student has already demonstrated they are unwilling

to or unable to learn the basics by cheating on them to begin with. Or Second, they must keep plagiarizing. This option ultimately leads to a student who does not know how to program potentially graduating with a Computer Science degree. If this student attempts to use this degree to apply for a job, the interview process, or, even worst, the opening weeks of a job, will reveal this lack of knowledge. Such a revelation is harmful to the entire community of computer science as it erodes confidence in the degree, and raises the bar for "entry level" job experience.

1.2. Prevalence

One may ask, "How prevalent is plagiarism in intro CS classes?" So did we, In an single 60 student intro level class during 2020, we observed an assignment where 10 people turned in exactly the same solution, and nearly 1/3rd of the class had easily recognizable plagiarized solution. This assignment will be referred to later as the "Motivating Case." This falls within the range of what other studies have found where between 1/3rd and 2/3rds of students have plagiarized for a assignment(Walcott, 2016)(Sheikh, 2004). Additionally, we have known many high level computer science students who don't seem to understand basic programming. While this lack of understanding may not point to plagiarism, that seems a reasonable culprit.

2. Definitions

Plagiarism is generally defined as "The practice of taking someone else's work or ideas and passing them off as one's own."(press, n.d.) However, for this paper we will define it exclusivity as "copying someone else's code to complete an assignment." We exclude copying ideas for several reasons, mainly because it is incredibly difficult to determine if an idea has been copied or just independently rediscovered; the general culture of computer science is to iterate on others ideas (often without attribution(Gibson, 2009)); and because in the

intro level, all the required ideas have been thought of by so many people that determining an origin would be next to impossible. Additionally, we exclude the possibility a student could “cite” their copied code to avoid plagiarism because the intro level is attempting to teach coding, not copying existing code. This definition works well for the intro level, however in higher levels it becomes less useful. We will use cheating and plagiarism synonymous in this paper. We would also like to define some theoretical plagiarism “vectors.”

2.1. Vectors

P2P Peer to Peer file sharing seems the most obvious vector. This is where one student writes the code and shares it with a peer who then submits the copy. In this instance one student is the author, and the other is the plagiarizer. Although both students are at fault, the plagiarizer is more so than the author.

Collaboration Collaboration is a potential vector for individual assignments where two students work on the assignment together even though it is intended to be a solo exercise. Collaboration is a particularly tricky vector as some level of collaboration is generally acceptable. This vector is distinct from P2P because neither student can be pointed to as a distinct “source”.

Theft While technically a subvector of P2P, Theft warrants its own definition. Theft is any form of P2P plagiarism where the author does not know that it is happening. For example, a student (the author) works on a shared computer and leaves a file saved locally, and the plagiarizer finds the file when attempting the same assignment.

Search Search engines, such as Google, are powerful tools, one of their powers is the ability to take an arbitrary string (Say, the title of a homework assignment for example) and get a useful result (like say, a github repo with the solution in 5 different programming languages). We are defining any time a plagiarizer uses an internet search engine to find an existing solution to an assignment to be this vector of plagiarism.

Expert Expert sources also exist. These generally fall into 4 categories.

- Paywalled answer repositories: This includes sites like Chegg, or Coursehero. On these (and similar sites), students can upload assignments to be done by expert programmers for a fee, or peruse other expert solutions hidden behind a paywall.

- Freelance Coders: This includes any professional coder who will write code for student assignments.
- Unethical Tutors: This includes any person who purports to be a tutor, but completes assignments directly instead of guiding their student.
- Large Language Models: This includes openly available generative models that can write code such as ChatGPT or Google’s Bard. These have gotten to the point where they can generate a bespoke solution to almost all intro level assignments if given an assignment as a prompt.

All variants of Expert are distinct from the Search vector for 2 reasons, first it is generally not free (or at least not freely accessible), and second the solution is tailored by the expert to the exact assignment given. This renders any attempt to “Scrape” answers futile. This appears to have been the primary vector in the Motivating Case, the 10 identical submissions were all sourced from a Chegg “Expert Answers” pay-walled solution repository.

It should be noted that all of these vectors can affect anything from single lines of code to entire assignments, and multiple vectors might be present in a single assignment. This makes it difficult to detect when they are happening, not only because it involves detecting plagiarism, but also because students are encouraged to use many of these vectors to aid their own work. As such we will be focused more on substantial cases where a potential plagiarizer acquires a significant amount of an assignment from the vector. For example, if a single line of code or a non required helper function is copied, probably not a problem; but if an entire required function is copied, that is a problem.

3. Current Approach

There are many existing research projects attempting to detect Plagiarism. The general method of plagiarism can be summarized as follows. Given a set A of submissions find a set P such that $\forall p \in A$, if $\exists q \in A$ such that $q \neq p$ and $D(p, q)$, then $p \in P$ and $q \in P$ for some plagiarism detection function D . Defining D effectively can be tricky (Walcott, 2016), however there have been several good attempts.

Researchers at Stanford University developed a program called MOSS (Measure of Software Similarity) which is capable of giving each assignment pair a similarity score. This approach is robust against name changes, code reorganization, and white space changes by analyzing a semi random subset of k-gram hashes based on each assignment. k has been manually tuned

for each language(Schleimer, Wilkerson, & Aiken, 2003).

Some have attempted to define D by ignoring syntax and exclusively using code semantics(Konecki, Orehovacki, & Lovrencic, 2009) so that 2 pieces of code match when they functionally do the same thing. We disagree with this specific approach for intro level classes, as many intro level problems should be have semantically identical solutions.

There are also many other efforts to define D using a machine learning model(Heres, 2017)(Katta, 2018)(Ullah, Wang, Farhan, Habib, & Khalid, 2018)(Awale, Pandey, Dulal, & Timsina, 2020).

Unfortunately, regardless of the definition used for D , all of these models have similar shortcomings when attempting to detect a given vector.

3.1. Theoretical Per Vector Performance

3.1.1. P2P. P2P is handled fairly well by these models. This vector is well suited for a source code analysis so long as D is resistant to minor changes and all sources used are submitted by someone. However, these models can not detect the author specifically.

3.1.2. Collaboration. Similarly to P2P, Collaboration is often detected by these models as the source code is quite similar. D however is harder to design for these as the “minor” changes can become more extreme with collaboration than with copying. Additionally, when they do detect plagiarism, these models can not discriminate between P2P and Collaboration.

3.1.3. Theft. Theft is as easy to detect as P2P and these models detect it well, however they can not discriminate between the plagiarizer and the victim, nor can they discriminate between P2P and Theft.

3.1.4. Search. Search is where existing models begin to substantially fail. No matter how sophisticated D is, it requires all potential sources to be present for consideration. Given the vast number of potential sources that can be found on the internet, detecting this vector is extremely difficult for D and generally requires a human to put in the leg work to scrape together the list of potential sources. Additionally, these searches can take $O(s^2 + se)$ Time (where s is student submissions, and e is external examples), for large values of e , this can be prohibitively long.

3.1.5. Expert. Expert is where existing models unambiguously fail the most. A properly executed Expert vector is indistinguishable from a legitimate attempt for all possibly D . As this content is Bespoke or behind a Paywall, it is genuinely impossible to provide the model with all relevant sources. The only way these models can detect an Expert vector is if 2 students both submit the same Expert source by accident.

3.2. General Shortcomings

In addition to the per vector performance, all of these models suffer from frequent false positives especially in the intro level. Consider Hello World in java. There is effectively 1 way to write hello world in java, you have your class, your main, and your print. This is similar for nearly all other languages. Since nearly all intro classes start with hello world, if the first assignment is checked with the above methods, it is likely that every student will be marked as plagiarized, even if they all attempted the assignment properly and individually.

3.3. Notable Alternatives

There are other existing models that can potentially address these problems. One research group used machine learning to detect stylistic variation in essays(AlSallal, Iqbal, Palade, Amin, & Chang, 2019). This approach has potential, especially when applied across a student’s entire code corpus, and might be able to detect a change of Author compared to other assignments. However, this vector is still weak to a serial plagiarist using a single source consistently. Additionally our target is intro level students who may not yet have a defined “style.” We believe it would be reasonable to suppose an intro level student’s style would drift significantly during their first few courses. Additionally, we did not find anyone attempting a similar single student stylistic analysis approach within code, just within human language. As such, we did not consider using this approach.

There are also hyper-focused detection models that attempt to exclusively a detect single LLMs output.(Koike, Kaneko, & Okazaki, 2024) However, these are not particularly reliable at present(Kirchner, Ahmad, Aaronson, & Leike, 2023) and have a high false positive rate(Elkhatat, Elsaid, & Almeer, 2023). At the time we started this research, LLMs were not a vector. Additionally the hyper-focused nature, rapid change in both detectors and LLMs, the unreliability within detectors, and the fact that they will not detect any vectors outside of an LLM expert, has deterred us from consider using this approach.

4. Our Approach

4.1. Single Source Plagiarism Detection

For our approach, we decided that it is more useful to be able to determine if an individual file itself is plagiarized, rather than a collection. To enable this, we created a unique IDE based on the Processing IDE. When the IDE first launches it creates a persistent UUID file on the machine that installed it. This UUID acts as a machine or user ID, and is referred to as the InstallID and is assumed to be constant throughout the class. Additionally, whenever a project is created with the IDE, a second UUID is created unique to that project, we refer to this as the ProjectID. Both UUIDs are saved in a special hidden metadata comment in the program source file, we will refer to this as the metaComment.

Whenever a file is opened, the IDE checks the InstallID of the file and machine. If they differ, the IDE notes the new InstallID in the metaComment in an ordered list called the InfectionStack. A similar stack is also used to track paste events. As these stacks are effectively one stack with a way to discriminate between event types, we will use InfectionStack to refer to both.

Additionally, when any part of the code is copied zero width spaces (U+200B) are interspersed with the normal letters to encode as much data as possible. The exact amount of data varies based on the size of the copy, but it can include several copies of the InstallID, ProjectID, and InfectionStack. When pasted into a project, the encoded data is compared to that of the project receiving the code, and any mismatch UUID's are added to the InfectionStack. This encoding survives being sent over most messaging software such as Discord, and is generally resilient to partial pastes. Any paste without encoded data is logged as originating outside the IDE. Regrettably all email clients we tested strip out the zero-width spaces, as such emailed code copies register as originating outside the IDE.

In addition to copy tracking, the IDE also logs the time, location, and content of all edit events including copy, paste, cut, delete, and typing. Any mismatched UUIDs are included in the log for the respective paste event. This data is also included in the metaComment. Using this keylog, it is possible to reconstruct the entire coding process and trace all code pastes.

Lastly, if the IDE loads a file without a metaComment, it notes this as the first edit and logs the initial state of the file.

While it can be enough to use a single file to detect Plagiarism, the UUIDs can be used to trace code authorship through multiple student submissions. Additionally, analyzing an entire class at the same time

can be useful to identify false positives such as copies from previous assignments on the student's machine.

4.2. Theoretical Per Vector Performance

This method was designed with our vectors in mind, and performs well on all of them.

4.2.1. P2P. As with existing methods, P2P is easy to detect. There is no easy way to submit someone else's unmodified file without the UUIDs being a clear red flag. Any attempt to edit the file will mark it as infected. Intro classes frequently have students add an "Author" comment to the top of their files. This would be needed to be edited by anyone attempting to cheat, and the original comment is preserved in the log. Finally any code shared via online messaging will be flagged at the very least, and generally, traced. If a student copies from a previous assignment, it will be traceable to which, and who authored that assignment.

Even simple files such as Hello World would be detected, as the detection does not rely on code comparison. Additionally, our method is immune to traditional false positives as organically created code is never considered plagiarized regardless of similarity to another student.

4.2.2. Collaboration. Collaborating students will have similar code with overlapping timestamps at each edit. Additionally, collaborating students tend to send code back and forth, meaning both sides' InfectionStack will be populated with each other's UUIDs. This back and forth should make it easy to distinguish from P2P.

4.2.3. Theft. Theft is detected in the same way as P2P. Unfortunately it can be tricky to discriminate it from P2P, but it is now theoretically possible to isolate the original file using meta knowledge of the computer setup (i.e. class computers would have a known InstallID), the InfectionStack and the timestamp log.

4.2.4. Search. All Search vectors are soundly defeated. No matter how a file is found, once it is loaded by the IDE, it will be marked as external. If it is submitted without being opened by the IDE, it will lack a metaComment. If the source code is copied into an existing project, it will be marked as external. The only viable source that this does not detect is physically typing out the code from a reference. This is theoretically detected by analyzing the timestamp data. Organically written code has a lot of back and forth. For example, Java's Hello World tends to be typed as

```

class HW{
}
[UP ARROW]
    public static void main(String[] args){
    }
[UP ARROW]
    System.out.println("Hello World");
[RUN]

```

Where as a plagiarist copying it by typing would likely write it top to bottom without backtracking significantly.

4.2.5. Expert. Expert sources are detected in the same way as Search. The only way an Expert can avoid detection as an external source is by installing the IDE and using that, in which case *they* will have a different InstallID from the student. This may require multiple assignments to detect, but is likely reliable over a full class unless the student finds 1 consistent Expert, or the Expert uses the student's own computer.

4.3. Evasion

There are notable ways to avoid detection by this method. In general we have dismissed them as "being more work than just *doing* the assignment." The metaComment is a plain text java comment appended to the file. Any text editor *other* than our IDE will show this comment and allow it to be edited. Convincingly faking the edit log would be incredible difficult, and deleting the entire string is easily detected.

A more viable method would be to share the InstallID file with an Expert, however, you would have to know that the InstallID exists, and where it is saved.

A corrupted metaComment is another possibility. The comment could be tampered with in such a way that it is incomplete. Sadly, this occasionally legitimately does happen if saving is interrupted so this is a concern.

We have dismissed these weaknesses reasoning that if an intro level student discovers *and* is able to exploit them, they likely know enough to not *need* to plagiarize in the intro level. And as the program is currently not well known, there will be no external assistance in doing so.

A potential bypass is to complete the assignment via any means that doesn't add to the InfectionStack, and then copy and paste the file content into another copy of the IDE, this then gets logged as an internal paste event between assignments and all keystroke logging is lost. This method we dismiss as there is little reason to do it (other than IDE bugs) and it is easily detectable as an event of interest.

4.4. Automated Detection

It is all well and good to wax eloquently about the theoretical ways the data can be used to detect plagiarism, and quite another to do so in an automated fashion suited to someone grading dozens of assignments. As such, we also developed a script that takes all student submissions for a given assignment and analyzes them. This script was developed after manual analysis of vectors within the Case Study, and as such, we did not implement detection of the features not present in the case study. Notably, we did not implement the timestamp based Collaboration detection as there was no Collaboration that would require such detection. We did not implement a way to discriminate P2P and Theft as no Theft occurred. Additionally we did not implement the linear coding detection as only 1 student used it and there were not enough students who had not plagiarized to define a heuristic for how linear is too linear.

The Script will take in a directory of submissions. First it will build a graph using the InfectionStack in an attempt to trace shared files. At the same time it builds a list of shared machines, which is used later. Any connections discovered at this point will be noted. Next the script goes student by student and checks all paste events. Purely internal paste events, or ones that are too short to track, are ignored. From there pastes fall into 4 categories: pastes originating from a project that was turned in by another student, or pastes originating from a different student's machine; pastes originating from the same machine, but a different project; pastes that originate from an InstallID that is not associated with any project turned in; and pastes of external origin. In each case where plagiarism is detected, the code pasted is also logged for human validation.

In some cases we use the number of lines pasted to determine if an event is plagiarism. This relies on an arbitrary threshold, however the authors of MOSS assert that there is a "sharp threshold" for a similar parameter within their algorithm (Schleimer et al., 2003). Unfortunately, this parameter is also domain specific and this threshold value is not given. We have assumed that this means picking a threshold is fairly intuitive and have picked the thresholds as seemed reasonable to us.

Pastes from another turned in ProjectID or InstallID that is associated with another student is immediately marked as plagiarism.

Pastes that originate on the same machine are treated differently depending on the machine. If the machine was not shared, the paste is assumed to originate from the same student. However there is the possibility they

are trying to use the loophole where the assignment is completed in another project then pasted to the final project to clean up the keystroke log. To detect this we decided to use a 50 line threshold. This should allow functions to be copied from previous assignments, but still detect entire projects being copied. If the machine was shared, the paste is flagged as plagiarism.

Pastes that originate from a different machine that isn't associated with any student and a project not associated with any student are assumed to be from the same student similar to pastes from the same machine, however, as this is less likely to be the student and may be an expert source, we have decreased the threshold to 20 lines.

Pastes that did not originate in the IDE are treated as copying from online resources. While a certain amount of online copying is permissible (like copying the name of a function from a docs page), copying whole functions would not be. We set the threshold for these events at 3 lines reasoning that 1 line is likely a line from a docs page, and 3 lines is likely accidentally copying a newline before and after the function within the docs page, but more than 3 lines could be an entire function.

Additionally, we track the number of edits made after the last plagiaristic paste event. We observe some students who initially copy a solution, and then spend considerable time adapting the solution (see 5.1E and 5.2F). It is not clear how such a case should be treated, so this number is recorded to allow a grader better context in such cases.

5. Case Study

During the Spring 2023 ICS 111 and 211 Intro to Java classes, we offered an extra credit assignment to students. Students were instructed to use the IDE and were expressly allowed to plagiarize the assignment. We worked with the professors of the classes to give each class a project that would be slightly too difficult to the average person in the class to encourage plagiarism, but not so difficult that NO students in the class would be able to complete it on their own. The 111 class was given a project to make 3 Polymorphic Shapes objects that move around the screen and bounce off the edges. The 211 class was given the task to create any function defined fractal (such as a Mandelbrot fractal). ChatGPT was able to generate solutions to both with just the prompt and no other assistance, and a solution for each was seeded on to an assortment of websites. Assignments were collected completely anonymously by distributing and collecting identical thumbdrives containing the IDE which was set to save to the drive its self. Assignment credit was given for

turning in a thumbdrive without checking its content. 2 students in 211 caught this loophole and did in-fact turn in completely blank thumbdrives.

Students were also asked to self report the method they used to complete the assignment and include the report on the thumbdrive. As students were asked to “cheat” on this assignment, the frequency of cheating its self carries no meaningful information, it only serves to reveal the Vectors and usability of our method.

As the 111 and 211 classes are separate assignments and cohorts, we will treat them as 2 concurrent Case Studies.

As part of each Case Study, the metaComment was stripped out of the files, and they were given to Stanford MOSS.

The code, raw data, pre-processed data, automation reports, and MOSS reports are available on Github¹

5.1. ICS 111

We received 6 submissions from 111, these have each been arbitrarily given a single letter designation between A and F. The results are summarized in Table 1. Of these 6, 4 successfully completed the assignment by some means, and 2 (A and D) did not.

Student A copy and pasted from ChatGPT, but prompted it wrong, and did a similar, but different assignment. Our metaComment makes it easy to see this, there is a 75 line paste followed by about 38 minor tweaks and formatting changes.

Student D appears to have struggled: the code they copied is a badly mangled AWT example and does not compile without significant modification.

Of the correct assignments, Student B completed the assignment legitimately and it is shown in the metaComments. There are multiple well organized files each organically coded with no large external pastes or irregularly linear typing sections. There are copies between files where code should be reused and modified, and follow up edits making these modifications.

Student C used the cross copy paste exploit mentioned previously and did not leave a comment. As such, we can not determine if their code is legitimate or not.

Student E pasted a large section of code from a different project with a matching InstallID and then spent considerable time (around 8 hours spread over a week or so) debugging and improving it. We believe this to be a legitimate completion.

Student F found and submitted the one of the seeded assignment file. It is logged as a large foreign paste.

¹<https://github.com/Jesse-McDonald/PlagerismIDE/tree/main/data>

Table 1. ICS 111 Case study

Student	Method	Visible in metaComment	Automated Check output	Moss Output
A	Plagiarized	Records large paste	Plagiarism Detected, 35 Edits	No Plagiarism
B	Legitimate	No Warning Signs	No Plagiarism Detected	No Plagiarism
C	Unclear	Large internal paste	Likely Plagiarized, 31 Edits	Minor link to E
D	Plagiarized	Records large paste	Plagiarism Detected, 31 Edits	No Plagiarism
E	Legitimate	Large Paste followed by many edits	Plagiarism Detected, 767 Edits	Minor link to C
F	Plagiarized	Large external paste	Plagiarism Detected, 0 Edits	No Plagiarism

5.1.1. MOSS. MOSS only identified a small connection between Student C and E. This connection is entirely composed of expected driver code.

5.1.2. Observed Vectors. Between all the submissions metaComment we were able to identify that Student A and F used Expert or Search vectors; Student B and E completed the assignment on their own; Student D used Search vector; and only student C is of indeterminate origin, but has a clean InfectionStack.

5.2. ICS 211

We received 18* submissions from 211, these have each been arbitrary given a single letter designation between A and R. The results are summarized in Table 2. Of these 18, 11 successfully completed the assignment by some means, 5 (C, G, L, M, and N) did not, and 2 (J and K) were completely blank thumbdrives.

Student C managed to find and copy a working Mandelbrot fractal written using Swing, which does not run using Processing.

Student G found and copied code for a Koch Snowflake written in Processing. Unfortunately, they found a badly written Koch Snowflake that doesn't actually run. We managed to fix their solution so that it would run, but when it did, it did not create any fractal.

We are unsure how C and G managed to find broken solutions before finding working solutions as we have been unable to find the original source code, but did find many working solutions while looking.

Student L found and copied a swing that makes a 3d render of a rain of rotating toroids. They even managed to convert it to run in Processing. However, this is not a solution to the actual the assignment given.

Student M used the cross copy paste exploit mentioned previously and did not leave a comment explaining how they completed the assignment. They did leave a comment explaining that it was an IDE stalling on saving bug that made them do the cross copy. This may indicate a large edit history was being saved, which may indicate a legitimate completion, but

we can not be sure this is the case. The solution also does not work, but with minor tweaking it can make 3 different recursive fractals, an infinite bulls eye pattern, a line with circles at every intersection and a Tetra-circle fractal. Unfortunately, the assignment specifically called for non-recursive fractals, so this is not a valid solution.

Student N appears to have legitimately coded the start of a 3d solar system simulation using rotation and translation matrices to enable a mouse to change camera angle and zoom. The planets in it do not orbit, but otherwise it is quite impressive, especially for having no external code usages. However, it is not a solution to the given assignment.

Of the correct assignments, a group of 5 (E, F, H, I, and P) were flagged by MOSS as having quite similar assignments, so we shall examine those first.

Student E's solution contains no large external pastes, however when the edits are animated, the result is visibly linear coding as if typing. The student claims they watched a youtube video of someone code a Mandelbrot fractal and followed along. This makes it a good candidate to test Linear typing detection, but we did not implement such detection as this is the only positive example, and we have limited legitimate completions to compare. While we consider E's solution to be Plagiarized, It is not part of a larger group unless they copied the same video.

Student F and P independently copied a stock example from the Processing website authored by Daniel Shiffman².

Student H and I followed a video by Coding Train where the same code by Daniel Shiffman is transformed from a Mandelbrot fractal to a Julia Set. After which, the final code was sent to I. The metaComment for I is corrupted after this paste, but the final result is altered from this share on both sides. The InfectionStack successfully tracks this exchange.

Student A found and copied a Penrose Snowflake example using multiple copies with no meaningful modifications.

Student B found and copied a Mandelbrot Fractal with zoom and made no meaningful modifications.

²<https://processing.org/examples/mandelbrot.html>

Table 2. ICS 211 Case study

Student	Method	Visible in metaComment	Automated Check output	Moss Output
A	Plagiarized	Records multiple large pastes	Plagiarism Detected, 1 Edits	No Plagiarism
B	Plagiarized	Records multiple large pastes	Plagiarism Detected, 0 Edits	Minor link to R
C	Plagiarized	Records multiple large pastes	Plagiarism Detected, 4 Edits	Minor link to R
D	Plagiarized	Records large paste	Plagiarism Detected, 0 Edits	Minor link to Q
E	Plagiarized	Visible Linear Coding	No Plagiarism Detected	Major Link to F, I, H, and P
F	Plagiarized	Records large paste	Plagiarism Detected, 105 Edits	Major Link to E, I, H, and P
G	Plagiarized	Records large paste	Plagiarism Detected, 78 Edits	
H	Plagiarized	Records large paste	Plagiarism Detected, 78 Edits	Major Link to E, F, I, P
I	Plagiarized	Infected by H InstallID	Copied Directly from H	Major Link to E, F, H, P
J	Blank	Blank Submission		
K	Blank	Blank Submission		
L	Plagiarized	Records multiple large pastes	Plagiarism Detected, 11 Edits	No Plagiarism
M	Unclear	Records Large Internal Paste	Plagiarism Detected, 0 Edits	No Plagiarism
N	Legitimate	No Warning Signs	No Plagiarism Detected	No Plagiarism
O	Plagiarized	Records multiple large pastes	Plagiarism Detected, 303 Edits	No Plagiarism
P	Plagiarized	Records large paste	Plagiarism Detected, 38 Edits	Major Link to E, F, I, H
Q	Plagiarized	Records large paste	Plagiarism Detected, 3 Edits	Minor link to D and R
R	Legitimate	No Warning Sign	No Plagiarism Detected	Minor Link to Q

Students D and Q both asked ChatGPT to generate a Mandelbrot, then copied the result without modification. The results are nevertheless distinct from each other.

Student O found and copied a tree fractal example, then modified it to be slightly different.

Student R claims to have used a video to help them make a Mandelbrot fractal complete with zoom. However there is no indication they copied from the video, and we believe this to be a legitimate completion of the assignment, all be it with a reference.

5.2.1. MOSS. MOSS identified a weak connection between B and R, as well as C and R that are incidental.

Additionally, it identified a moderate connection between D and Q, as well as R and Q. The same connection is found despite the connection not being found between D and R, and only 2 of them being generated by ChatGPT. The marked code is also quite incidental (nested for loops with same iterator names).

MOSS Identified a strong connection between H and I, as well as between F and P. This connection is a successful detection of plagiarism within those 2 pairs. However, Moss also identifies weaker, but still strong connections between all 4, and additionally E. It may not be accurate to include E in this group, as E used a different source. The code between the other 4 though has the same original author, with modifications. It is tempting to say that MOSS correctly identified the plagiarism despite the changes, however the similarities between the 2 groups are the same as the similarities between each and E, and the code is quite common to

Mandelbrot fractals. As such, we believe this to be incorrect with regard to plagiarism.

5.2.2. Observed Vectors. Between all the submissions metaComment we were able to identify that Students A-C, F-I, L,O, and P used the Search Vector; Students N and R completed the assignment on their own; Students D and Q used Expert vector; Student E used search but masked it with linear typing; Students J and K exploited loopholes in grading; and only student M is of indeterminate origin, but has a clean InfectionStack.

6. Conclusion

In this work, we have assessed the potential merit of edit event logging metadata generated by a specialized IDE for tracking plagiarism. We find that this approach works quite well even against formerly impossible vectors such as Expert sources and undefinable Search results that current approaches struggle to detect. This emphasizes the weakness of current approaches that compare two code samples to determine similarity as a way to detect plagiarisms, it is not possible to provide all the samples such a system requires, and the existence of code generating AI models only exacerbate the existing problem.

Our approach has shown promise as a tool in detecting plagiarism within intro level CS courses. The most useful features are paste event tracking and UUID tracking. We encourage anyone concerned with

plagiarism in an intro level class to provide a modified IDE for their class that uses these or similar tracking features.

In the future, we hope to completely automate the detection of plagiarized assignments created with the IDE including linear typing, and timestamp based collaboration detection.

Additionally, our primary aim is to research plagiarism reduction methods, so it may be equally useful to create a heuristic that represents the general amount of plagiarism present in a class as a whole rather than individual assignments.

References

- AlSallal, M., Iqbal, R., Palade, V., Amin, S., & Chang, V. (2019, Jul). An integrated approach for intrinsic plagiarism detection. *Future Generation Computer Systems*, 96, 700–712. doi: 10.1016/j.future.2017.11.023
- Awale, N., Pandey, M., Dulal, A., & Timsina, B. (2020, Jul). Plagiarism detection in programming assignments using machine learning. *September 2020*, 2(3), 177–184. doi: 10.36548/jaicn.2020.3.005
- Elkhatat, A. M., Elsaid, K., & Almeer, S. (2023, Sep). Evaluating the efficacy of ai content detection tools in differentiating between human and ai-generated text. *International Journal for Educational Integrity*, 19(1). doi: 10.1007/s40979-023-00140-5
- Gibson, J. P. (2009, July). Software reuse and plagiarism: a code of practice. *ACM SIGCSE Bulletin*, 41(3), 55–59. Retrieved from <http://dx.doi.org/10.1145/1595496.1562900> doi: 10.1145/1595496.1562900
- Heres, D. (2017). *Source code plagiarism detection using machine learning* (Unpublished doctoral dissertation). Utrecht University.
- Katta, J. Y. B. (2018). *Machine learning for source-code plagiarism detection*.
- Kirchner, J. H., Ahmad, L., Aaronson, S., & Leike, J. (2023, Jan). *New ai classifier for indicating ai-written text*. Retrieved from <https://openai.com/blog/new-ai-classifier-for-indicating-ai-written-text>
- Koike, R., Kaneko, M., & Okazaki, N. (2024, March). Outfox: Llm-generated essay detection through in-context learning with adversarially generated examples. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(19). Retrieved from <http://dx.doi.org/10.1609/aaai.v38i19.30120> doi: 10.1609/aaai.v38i19.30120
- Konecki, M., Orehovacki, T., & Lovrencic, A. (2009, June). Detecting computer code plagiarism in higher education. In *Proceedings of the iti 2009 31st international conference on information technology interfaces*. IEEE. Retrieved from <http://dx.doi.org/10.1109/ITI.2009.5196118> doi: 10.1109/iti.2009.5196118
- press, O. U. (Ed.). (n.d.). *Oxford english dictionary*.
- Schleimer, S., Wilkerson, D. S., & Aiken, A. (2003, June). Winnowing: local algorithms for document fingerprinting. In *Proceedings of the 2003 acm sigmod international conference on management of data*. ACM. Retrieved from <http://dx.doi.org/10.1145/872757.872770> doi: 10.1145/872757.872770
- Sheikh, L. (2004). Is plagiarism more prevalent in some forms of assessment than others.. Retrieved from <https://api.semanticscholar.org/CorpusID:210566353>
- Ullah, F., Wang, J., Farhan, M., Habib, M., & Khalid, S. (2018, Oct). Software plagiarism detection in multiprogramming languages using machine learning approach. *Concurrency and Computation: Practice and Experience*, 33(4). doi: 10.1002/cpe.5000
- Walcott, P. (2016, December). Attitudes of second year computer science undergraduates toward plagiarism. *Caribbean Teaching Scholar*, 41(6), 63–80.